



A next-gen AI assistant for your PDFs — ask questions, get instant answers, and manage conversations like a pro!

About

The **Smart PDF Chatbot** is a **Generative AI-powered assistant** that transforms static PDF documents into an interactive knowledge base. Ask natural language questions, get **context-aware answers**, and interact with multiple PDFs without manual searching. Think of it as **ChatGPT meets your PDF library**.

Why This Project

- Tired of scrolling PDFs for key information?
- Want a **dynamic Q&A** experience with your documents?
- Need **persistent chat sessions** to continue your research?

This project solves all of that with:

- **RAG (Retrieval-Augmented Generation)**
- **FAISS vector storage**
- **HuggingFace embeddings**
- **Streamlit-based interactive UI**

Key Features

Feature	How It Helps You
Upload PDFs	Instantly process one or multiple PDFs
RAG-powered answers	Vector search + LLM reasoning
Chat sessions	Save, rename, and continue past conversations
Modern UI	Dark, glassy interface with smooth animations
Smart responses	If info isn't in the PDF, bot tells you politely
Dynamic updates	Add/update/remove PDFs seamlessly

How It Works

1. **Upload PDFs:** Documents loaded and parsed with `PyPDFLoader`.

2. **Text Chunking:** Large PDFs split into manageable text chunks.
 3. **Vector Embeddings:** HuggingFace model generates embeddings.
 4. **FAISS Storage:** Efficient vector search for retrieving relevant chunks.
 5. **LLM Response:** GPT, Gemini, or Groq generates contextual answers.
 6. **Session Management:** Each chat is stored, retrievable, and renameable.
-



Tech Stack

- **Frontend:** Streamlit
 - **LLM Integration:** OpenAI / Gemini / Groq
 - **Document Processing:** LangChain (`PyPDFLoader`, `RecursiveCharacterTextSplitter`)
 - **Embeddings:** HuggingFace Sentence Transformers
 - **Vector Storage:** FAISS
 - **State Management:** Streamlit `session_state`
 - **UI Design:** Custom CSS, dark/glassy theme
-



Project Layout

- `app.py` → Streamlit UI & main app
 - `rag_pipeline.py` → RAG logic (retrieval + prompt construction)
 - `vectorstore_manager.py` → FAISS & embeddings
 - `chat_OpenAI.py` / `chat_Gemini.py` → LLM wrapper
 - `session_manager.py` → Session handling
 - `history_manager.py` → Saves conversation history
 - `data/` → User-uploaded PDFs
 - `requirements.txt` → Python dependencies
-



Quick Start



Clone & Navigate

```
git clone https://github.com/aruk3002/RAG_PDF_CHATBOT.git  
cd smart-pdf-chatbot
```



Setup Virtual Environment

```
python -m venv venv  
source venv/bin/activate # Mac/Linux  
venv\Scripts\activate # Windows
```

Install Dependencies

```
pip install -r requirements.txt
```

Add API Keys

Create a `.env` file in project root:

```
OPENAI_API_KEY=your_api_key  
GEMINI_API_KEY=your_api_key  
GROQ_API_KEY=your_api_key
```

Run the App

```
streamlit run app.py
```

Open your browser: <http://localhost:8501/>

Example Conversation

User:

What does manufacturing involve?

Bot:

Manufacturing is not explicitly defined in the document, but it discusses processes like production efficiency, machinery, and quality control. Essentially, it converts raw materials into finished products.

Updating Knowledge Base

- **Add PDFs:** Upload → embeddings & FAISS updated automatically.
- **Replace PDFs:** Re-upload → re-embedding for only updated files.
- **Remove PDFs:** Delete → reflected in next session.

No retraining required!

Future Enhancements

-  Keyword-based suggested questions
 -  Persistent FAISS index between sessions
 - Multi-user session isolation
 -  Chat export & delete options
-

Contributing

1. Fork the repo
 2. Create a feature branch
 3. Submit a pull request
-



Chandan Aruk



Data Analyst |  AI & Data Science Enthusiast

[LinkedIn](#)